

Sequential Monte-Carlo Filtering/Smoothing for State and Parameter Estimation in General Nonlinear State Space Systems

Ali Gorji

May 30, 2018

- Problem statement: **Nonlinear State Space Systems (NSSS)**
- State and parameter estimation in NSSS
- Expectation Maximization (EM) for joint state and parameter estimation in NSSS
- Handling non-linearity: Sequential Monte Carlo (SMC)
- Particle Filter/Smoother for joint state and parameter estimation

Nonlinear State Space Systems

- We consider a family of **discrete-time** and **affine** NSSSs:

$$\begin{aligned}\mathbf{x}_k &= f_{\theta^x}(\mathbf{x}_{k-1}, \mathbf{u}_k) + \mathbf{v}_k \\ \mathbf{y}_k &= g_{\theta^y}(\mathbf{x}_k) + \mathbf{w}_k\end{aligned}$$

- Where:
 - \mathbf{x}_k : hidden state at the k -th step-time
 - \mathbf{y}_k : output observation at the k -th step-time
 - f, g : parametric state transition and output models with parameters θ^x and θ^y , respectively
 - \mathbf{u}_k : external input
 - $\mathbf{v}_k, \mathbf{w}_k$: additive noise

NSSS in Time-series Modeling with Missing Data

- Considering \mathbf{y}_k as a vector of output observations at the k-th step:
 - Time-series may suffer from **missing-data** where some of outputs are not **measurable**
 - **Weather-data**: daily max/min temperature is available, but the associated time-of-year may be missing
- **NSSS**: model time-series as a state-space system with missing-data being represented by **hidden-states**
- Training the model, **infer** missing-data from the received observations

NSSS in Model Reference Control

- The goal in Model Reference Control (MRC):

$$\mathbf{u}_{1:K} = \underset{\mathbf{u}}{\operatorname{argmin}} \left\{ \sum_k (y_k - \hat{y}_k)^2 \right\}$$

with \hat{y}_k being a desired trajectory

- Consider a parametric structure for the controller: $\mathbf{u}_k = h_{\theta^u}(\mathbf{x}_{k-1})$ and re-construct the NSSS:

$$\mathbf{x}_k = f_{\theta^x}(\mathbf{x}_{k-1}, h_{\theta^u}(\mathbf{x}_{k-1})) + \mathbf{v}_k$$

- Jointly estimate states and parameters to minimize MRC objective
- Generate control command using the learned h_{θ^u}

NSSS in Observer Trajectory Planning

- Consider a moving object with motion state \mathbf{x}_k , and a moving observer with state \mathbf{x}_k^o that gathers measurements y_k
- NSSS in dynamical target tracking:

$$\mathbf{x}_k = A\mathbf{x}_{k-1} + B\mathbf{v}_k$$

$$y_k = g_{\mathbf{x}_k^o}(\mathbf{x}_k) + \mathbf{w}_k$$

with A, B being known matrices of state-transition model, and g as a nonlinear measurement model such as bearing

- The goal is to estimate observer locations to minimize **variance of target state estimation**

State and Parameter Estimation in NSSS

- Three main sources in NSSS:
 - Observations: $\{\mathbf{y}_{0:k}, \mathbf{u}_{0:k}\}$
 - Parameters: $\{\theta^x, \theta^y\}$
 - Hidden states: $\mathbf{x}_{0:k}$
- State estimation: estimate posterior density of hidden states given observations:
 - Filtering: $p(\mathbf{x}_{0:k} | \mathbf{y}_{0:k}, \mathbf{u}_{0:k})$
 - Smoothing: $p(\mathbf{x}_{0:k'} | \mathbf{y}_{0:k}, \mathbf{u}_{0:k}), k' < k$
- Parameter estimation: obtain an estimate of model parameters by optimizing a proper objective such as likelihood

$$\theta^* = \operatorname{argmax}_{\theta} \{p(\mathbf{y}_{0:k}; \theta)\}$$

Maximum Likelihood Estimator

- Parameters can be obtained by maximizing the log-likelihood function

$$\theta^* = \operatorname{argmax}_{\theta} \left(J = \sum_k \log \{ p(\mathbf{y}_k; \theta) \} \right)$$

- Incomplete likelihood:** the likelihood cannot be calculated since hidden states are missing
- Main trick:** given the concavity of log-function and using Jensen inequality:

$$\log \left\{ \underbrace{p(\mathbf{y}_k; \theta)}_{E_q \left(\frac{p(\mathbf{y}, \mathbf{x}; \theta)}{q(\mathbf{x})} \right)} \right\} \geq E_{q(\mathbf{x})} \left(\underbrace{p(\mathbf{y}_k | \mathbf{x}_k; \theta)}_{\text{Complete Likelihood}} \right) - \mathcal{H}(\mathbf{x}_k)$$

with $q(\cdot)$ being a distribution over states and $\mathcal{H}(\cdot)$ as the entropy

Expectation Maximization (EM) Algorithm

- Instead of maximizing the likelihood, we aim to maximize the **lower-bound**
- EM algorithm consists of two main phases:
 - **E-phase**: assuming parameters of the model are fixed, minimize the bound-gap through minimizing the following **kullback leibler divergence**:

$$D_{KL}(q||p) = -E_q \left(\log \left(\frac{p(\mathbf{x}_k|\mathbf{y}_k; \theta)}{q(\mathbf{x}_k)} \right) \right)$$

- **M-phase**: given estimates of hidden-states, maximize the expectation over the complete likelihood:

$$\theta^* = \operatorname{argmax}_{\theta} E_q \left(\log (p(\mathbf{y}_k|\mathbf{x}_k; \theta)) \right)$$

Iterative State/Parameter Estimation Using EM

- Initialize model parameters, training iteration $n = 0$ and model parameters $\theta[0]$
- Do until convergence:
 - **E-phase**: by estimating $q(\mathbf{x}_k) = p(\mathbf{x}_k|\mathbf{y}_k; \theta[n]), k \in \{0, \dots, K\}$
 - **M-phase**: update parameters by maximizing $\sum_{k=0}^K E_q (\log (p(\mathbf{y}_k|\mathbf{x}_k; \theta)))$ with q being used from the E-phase
 - Calculate log-likelihood using estimated states and new parameter values
 - If no significant change in log-likelihood **break**, otherwise set $n \leftarrow n + 1$
- **Theorem**: Under **exact state estimation**, EM guarantees log-likelihood increases at each step: $\mathcal{L}(\theta[n+1]) \geq \mathcal{L}(\theta[n])$ with $\mathcal{L}(\theta)$ being the complete likelihood

Nonlinearity and State/parameter Estimation

- Under additive Gaussian noises and linear state-transition and output models:
 - Optimal filtering such as **Kalman filter** and Gaussian distributed posterior $p(\mathbf{x}_k|\mathbf{y}_k; \theta) = \mathcal{N}(\mathbf{x}_{k|k}, P_{k|k})$ with $\mathbf{x}_{k|k}$ and $P_{k|k}$ being mean and covariance of estimates, respectively
 - Obtain a **closed-form** for $\sum_{k=1}^K E_q(\log(p(\mathbf{y}_k|\mathbf{x}_k; \theta)))$ and find the optimal parameter value
- Under general nonlinear state-transition and/or output models:
 - Local linearization and **Extended Kalman Filter** for approximate E-phase and Gradient updates for M-phase
 - Approximate the posterior distribution of states using samples and find an unbiased estimate of the log-likelihood that could be maximized against parameters of the model

Approximate State Estimation Using Sequential Monte Carlo (SMC)

- Main idea is to represent the posterior distribution as a weighted summation of samples:

$$p(\mathbf{x}_k | \mathbf{y}_{1:k}) \sim \sum_{m=1}^M \tilde{w}_k^m \delta(\mathbf{x}_k - \mathbf{x}_k^m)$$

with \tilde{w}_k^m being normalized **importance weights**, $\delta(\cdot)$ as the delta-Dirac function and M as the total number of **particles**

- Approximate M-phase with SMC sampler:

$$\theta^* = \underset{\theta}{\operatorname{argmax}} \sum_{k=1}^K \sum_{m=1}^M \tilde{w}_k^m \log(\mathbf{y}_k | \mathbf{x}_k^m; \theta)$$

Sequential Importance Sampling

- Draw samples from a **proposal distribution** $s(\mathbf{x}_{0:k}|\mathbf{y}_{0:k})$ and update importance weights recursively:

$$w_k^m = w_{k-1}^m \frac{p(\mathbf{x}_k^m|\mathbf{x}_{k-1}^m)p(\mathbf{y}_k|\mathbf{x}_k^m)}{s(\mathbf{x}_k^m|\mathbf{x}_{k-1}^m, \mathbf{y}_k)}$$

and normalize the weights $\tilde{w}_k^m = \frac{w_k^m}{\sum_{m=1}^M w_k^m}$

- Choice of proposal density function:
 - State-transition: $s(\mathbf{x}_k^m|\mathbf{x}_{k-1}^m, \mathbf{y}_k) = p(\mathbf{x}_k^m|\mathbf{x}_{k-1}^m)$
 - Using the outputs and sample from $p(\mathbf{x}_k^m|\mathbf{x}_{k-1}^m)p(\mathbf{y}_k|\mathbf{x}_k^m)$
 - Use sub-optimal Extended Kalman Filter (EKF) to obtain an estimate of the posterior $p(\mathbf{x}_k|\mathbf{y}_k)$ and assign the distribution to the proposal

Sequential Importance Re-sampling

- Importance sampling provides estimates whose variance increases **exponentially** over time
 - Only a few particles survive with a high weight with other particles will be assigned to a zero-weight
- We define **re-sampling** as an operation that maps the set $\{\tilde{w}_k^m, \mathbf{x}_k^m\}$ to a set of **equally-weighted** states $\{\frac{1}{M}, \bar{\mathbf{x}}_k^m\}$
- **Idea**: samples with higher weights will be replicated, but lower-weighted samples also survive
- Different re-sampling strategies:
 - Residual re-sampling, systematic re-sampling and multinomial re-sampling
 - Adaptive re-sampling: re-sample only if the **effective sample-size** of weights is below a threshold:

$$\left(\sum_{m=1}^M (\tilde{w}_k^m)^2 \right)^{-1} \leq \tau$$

with $\tau = \frac{M}{2}$

General SMC Sampler for Filtering

- Sample initial particles \mathbf{x}_0^m from an arbitrary distribution
- Compute initial normalized weights \tilde{w}_0^m
- Check the re-sampling criterion and if needed, run re-sampling and generate $\{\frac{1}{N}, \bar{\mathbf{x}}_0^m\}$
- For $k > 0$:
 - Draw new samples from the proposal distribution $s(\mathbf{x}_k^m | \bar{\mathbf{x}}_{k-1}^m, \mathbf{y}_k)$
 - Compute the normalized weights \tilde{w}_k^m
 - Check for the re-sampling criterion:
 - If effective sample-size is lower than the threshold, run re-sampling and $\{\tilde{w}_k^m, \mathbf{x}_k^m\} \rightarrow \{\frac{1}{M}, \bar{\mathbf{x}}_k^m\}$
 - If effective sample-size is above the threshold, use $\{\tilde{w}_k^m, \mathbf{x}_k^m\}$ as the set of particles and associated weights

Summary: EM with SMC

- **Initialization:** randomly initialize parameters $\theta[0]$ and M particles with associated weights $\{\frac{1}{M}, \mathbf{x}_0^m\}$
- Do until convergence:
 - **E-phase:** run one step of SMC sampler and obtain sampled states and associated weights $\{\tilde{w}_k^m, \bar{\mathbf{x}}_k^m\}, m \in \{1, \dots, M\}, k \in \{0, \dots, K\}$
 - **M-phase:** approximate the complete log-likelihood and obtain new parameters through $\theta[n+1] = \operatorname{argmax}_{\theta} \sum_{k=1}^K \sum_{m=1}^M \tilde{w}_k^m \log(\mathbf{y}_k | \bar{\mathbf{x}}_k^m; \theta)$
 - Stop if the likelihood does not change significantly $|\mathcal{L}(\theta[n+1]) - \mathcal{L}(\theta[n])| \leq \epsilon$ or $n \geq n_{\max}$, otherwise $n \leftarrow n+1$

State Estimation: Smoothing

- Given a **batch of observations** ($\mathbf{y}_{0:k}$), estimate $p(\mathbf{x}_k | \mathbf{y}_{0:K})$ instead of filtering distribution:
 - Makes state estimation **more robust** to the process or observation noise
 - Allows the training procedure to have a better estimate of model parameters
- Different perspectives:
 - **Forward-backward**: Use estimated states at the filtering stage (forward path) to approximate smoothing distribution
 - **Two-filter smoothing**: Break down the smoothing distribution into two separate filtering stages and estimate each stage through a single filter
 - **Maximum a Posteriori smoothing**: Run the forward path and estimate states. Then, obtain smoothed estimates by maximizing the posterior distribution of hidden states over measurements (through Dynamic Programming)

Forward-backward Particle Smoothing

- Modeling the smoothing distribution as
$$p(\mathbf{x}_k | \mathbf{y}_{0:K}) \sim \sum_{m=1}^M \tilde{w}_{k|K}^m \delta(\mathbf{x} - \mathbf{x}_k^m)$$
- Updating weights using the following forward-backward rule:

$$\tilde{w}_{k|K}^m = \tilde{w}_k^m \left\{ \sum_{m=1}^M \tilde{w}_{k+1|K}^m \frac{p(\mathbf{x}_{k+1}^m | \mathbf{x}_k^m)}{\sum_{m'=1}^M \tilde{w}_k^{m'} p(\mathbf{x}_{k+1}^{m'} | \mathbf{x}_k^{m'})} \right\}$$

- Some considerations:** the computational complexity is of order M^3
- There are approximations to reduce to $M^2 \log(M)$ by partitioning particles
- EM Particle Smoothing:** calculate the complete likelihood using samples generated by smoothing distribution

Summary

- Family of **discrete-time** nonlinear state space systems was discussed
- An NSSS is fully characterized through **hidden states** and its **model parameters**
- EM algorithm is used to estimate hidden states and learn parameters of the model
- SMC sampler (particle filtering) aims to provide an **approximate low-variance estimate** of the posterior distribution of states
- EM with particle filtering/smoothing forms a universal framework for jointly estimating hidden states and learning parameters of a general nonlinear state space system

- A. Doucet, and A. M. Johansen, “A tutorial on particle filtering and smoothing: fifteen years later”, in Handbook of Nonlinear Filtering (eds. D. Crisan et B. Rozovsky), Oxford University Press, 2011.
- M. Klaas, et al. “Fast particle smoothing: If I had a million particles””, Proceedings of the 23rd International Conference on Machine Learning, Pittsburgh, PA, 2006.
- S. Roweis, and Z. Ghahramani, “An EM algorithm for identification of nonlinear dynamical systems”, Proceedings of the 1998 conference on Advances in neural information processing systems.

Questions

